

Team Playbook

BATTLE OF THE COASTS • 72-HOUR HACKATHON • SEPT 12–14

Everything your team needs going in: the eight tracks you can build in, how the 72 hours are structured, when you're scored, and how those scores add up to a placement.

- DURATION72 hours
- TRACKS8 — one winner named per track
- CHECK-INSEvery 12 hours — scored, same rubric as the final
- PLACEMENT SCOREFinal round score + average of your 5 check-in scores

JUDGED BYOpenAI • Google • AWS • Anthropic • Sequoia • Lovable • Lightspeed • Underscore VC • HOF Capital • CADA

1. EVENT AT A GLANCE

72 hours, six scored moments. Teams come in from schools on both coasts — every 12 hours, judges score your current build against the same rubric they'll use at the final, so consistent progress is worth as much as a strong finish.

East	Harvard • MIT • Columbia • BU • NYU
West	Stanford • Caltech • Berkeley • UCLA • USC

Coast is about where a team was recruited from, not how it's judged — East and West teams compete side by side within the same eight tracks below. There's no separate East-only or West-only prize.

2. TIMELINE

0h	12h	24h	36h	48h	60h	72h
Kickoff	Check-in 1	Check-in 2	Check-in 3	Check-in 4	Check-in 5	FINAL
Rubric published	Scored	Scored	Scored	Scored	Scored	Scored

3. TRACKS

Pick one track at registration and build in it for the full 72 hours. Every track is judged on the same four-category rubric, and each names its own winner.

Track	What it covers	Winners
Life Sciences	AI applied to drug discovery, diagnostics, genomics, or healthcare delivery.	1
Edge AI	Models and systems that run on-device — phones, wearables, robotics, embedded hardware.	1
Cloud AI	Infrastructure, tooling, and platforms for training or serving models at scale.	1
Inference	Speed, cost, or efficiency breakthroughs in running models in production.	1

Track	What it covers	Winners
AI Apps	End-user products built on top of existing models — the application layer.	1
Human Flourishing	AI for wellbeing, education, mental health, creativity, or human connection.	1
Deep Tech / Physical AI	Robotics and embodied AI — hardware-software systems in the physical world.	1
Defense	AI for national security, defense, or dual-use applications.	1

4. SUBMISSION CADENCE

You'll submit something every 12 hours — six submissions in total, and every one of them is scored the same way.

Check-ins — hours 12, 24, 36, 48, 60

Post a link to your current build (repo, deployed demo, or working doc — whatever's real at that hour) at [submission link — organizer to add]. A judge scores it 1–10 across the same four categories used at the final.

- Each check-in score counts — the five are averaged into one number that becomes half of your placement score.
- Judges score the build in front of them at that hour, not against the finished product it might become — a “9–10” on Innovation at hour 12 means a breakthrough concept, not a finished demo.
- A weak check-in isn't fatal on its own, but five weak ones will cost you more than one bad final demo would.

Final submission — hour 72

Your finished project, submitted at [submission link — organizer to add]: repo/demo link, a short writeup, and whatever you're presenting live or on video. Scored the same way, by the full judging panel for your track.

5. WHAT'S JUDGED, EVERY ROUND

Four categories, same weights, used at every scored round — checkpoints and final alike.

Category	Weight	Question
Innovation & Creativity	30%	How novel and creative is the solution?
Technical Implementation	25%	How well is the solution built?
Business Value & Impact	25%	How valuable is this solution?
Presentation & Communication	20%	How well is the idea communicated?

6. HOW A TRACK IS WON

Your project is only ever compared to other projects in your own track — coast has no bearing on judging.

$$\text{PLACEMENT SCORE} = \text{FINAL ROUND SCORE} + \text{AVERAGE OF YOUR 5 CHECK-IN SCORES}$$

Each of the eight tracks is scored and ranked separately. The highest placement score within your track wins that track's prize. A strong Defense submission has no bearing on who wins Life Sciences, and an East Coast team and a West Coast team in the same track are judged against each other exactly the same way.

7. FAQ

Do check-in scores really count toward my final result?

Yes. Each check-in is scored 1–10 across the same four categories as the final. The average of your five check-in scores is added to your final round score to produce your placement score.

What happens if I miss a check-in?

A missed check-in scores as a 0 for that round and will pull your average down — treat every 12-hour mark as a real deadline, not just the final one.

Can I switch tracks mid-event?

No — pick your track at registration and build in it for the full 72 hours, so every one of your six scores lands in the right pool.

Does it matter which coast I'm from?

No. East and West describe where teams were recruited from, not how they're scored — everyone in a track is judged the same way.